Chapter 10

- 10.1 Numerical integration (overview)
- 10.2 Distributional approximations (overview, more in Chapter 4 and 13)
- 10.3 Direct simulation and rejection sampling (overview)
- 10.4 Importance sampling (used in PSIS-LOO discussed later)
- 10.5 How many simulation draws are needed? (Ex 10.1 and 10.2)
- 10.6 Software (can be skipped)
- 10.7 Debugging (can be skipped)

Notation

- In this chapter, generic $p(\theta)$ is used instead of $p(\theta|y)$
- unnormalized distribution is denoted by $q(\cdot)$
- proposal distribution is denoted by $g(\cdot)$

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\rightarrow$ 0

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\to 0$

  - closest value to 1 is $\approx 1 \pm 2.2 \cdot 10^{-16}$
    - pbeta(0.5,241945,251527) $\to 1$

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\rightarrow$ 0
    - pbeta(0.5,241945,251527, lower.tail=FALSE)
      $\approx -1.15 \cdot 10^{-42}$
  - closest value to 1 is $\approx 1 \pm 2.2 \cdot 10^{-16}$
    - pbeta(0.5,241945,251527) $\rightarrow$ 1

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\rightarrow$ 0
    - pbeta(0.5,241945,251527, lower.tail=FALSE)
      $\approx -1.15 \cdot 10^{-42}$
  - closest value to 1 is $\approx 1 \pm 2.2 \cdot 10^{-16}$
    - pbeta(0.5,241945,251527) $\rightarrow$ 1
- Log densities
  - use log densities to avoid over- and underflows in floating point presentation
    - sum(dnorm(qr,log=TRUE)) $\rightarrow$ -847.3

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\to 0$
    - pbeta(0.5,241945,251527, lower.tail=FALSE)
      $\approx -1.15 \cdot 10^{-42}$
  - closest value to 1 is $\approx 1 \pm 2.2 \cdot 10^{-16}$
    - pbeta(0.5,241945,251527) $\to 1$
- Log densities
  - use log densities to avoid over- and underflows in floating point presentation
    - sum(dnorm(qr,log=TRUE)) $\to$ -847.3
    - how many observations we can now handle?

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\rightarrow$ 0
    - pbeta(0.5,241945,251527, lower.tail=FALSE)
      $\approx -1.15 \cdot 10^{-42}$
  - closest value to 1 is $\approx 1 \pm 2.2 \cdot 10^{-16}$
    - pbeta(0.5,241945,251527) $\rightarrow$ 1
- Log densities
  - use log densities to avoid over- and underflows in floating point presentation
    - sum(dnorm(qr,log=TRUE)) $\rightarrow$ -847.3
    - how many observations we can now handle?
  - compute exp as late as possible
    - e.g. for $a > b$, compute
      $\log(\exp(a) + \exp(b)) = a + \log(1 + \exp(b - a))$
      e.g. $\log(exp(800) + exp(800)) \rightarrow$ Inf, but
      $800 + \log(1 + exp(800 - 800)) \approx 800.69$

Numerical accuracy

- Floating point presentation of numbers. e.g. with 64bits
  - closest value to zero is $\approx 2.2 \cdot 10^{-308}$
    - qr=rnorm(600);prod(dnorm(qr)) $\rightarrow$ 0
    - pbeta(0.5,241945,251527, lower.tail=FALSE)
      $\approx -1.15 \cdot 10^{-42}$
  - closest value to 1 is $\approx 1 \pm 2.2 \cdot 10^{-16}$
    - pbeta(0.5,241945,251527) $\rightarrow$ 1
- Log densities
  - use log densities to avoid over- and underflows in floating point presentation
    - sum(dnorm(qr,log=TRUE)) $\rightarrow$ -847.3
    - how many observations we can now handle?
  - compute exp as late as possible
    - e.g. for $a > b$, compute
      $\log(\exp(a) + \exp(b)) = a + \log(1 + \exp(b - a))$
      e.g. $\log(exp(800) + exp(800)) \rightarrow$ Inf, but
      $800 + \log(1 + exp(800 - 800)) \approx 800.69$
    - e.g. in Metropolis-algorithm compute the log of ratio of densities using the identity
      $\log(a/b) = \log(a) - \log(b)$

It's all about expectations

$$E_\theta[f(\theta)] = \int f(\theta)p(\theta|y)d\theta$$

- Conjugate priors and analytic solutions
- Grid integration and other quadrature rules
- Independent Monte Carlo, rejection and importance sampling
- Markov Chain Monte Carlo
- Distributional approximations (Laplace, VB, EP)

Quadrature integration

- The simplest quadrature integration is grid integration
  - Evaluate function in a grid and compute

$$E[-\alpha/\beta] \approx \sum_{t=1}^{T} w_{\text{cell}}^{(t)} \frac{\alpha^{(t)}}{\beta^{(t)}},$$

where $w_{\text{cell}}^{(t)}$ is the normalized probability of a grid cell t, and $\alpha^{(t)}$ and $\beta^{(t)}$ are center locations of grid cells

Quadrature integration

- The simplest quadrature integration is grid integration
  - Evaluate function in a grid and compute

  $$E[-\alpha/\beta] \approx \sum_{t=1}^{T} w_{\text{cell}}^{(t)} \frac{\alpha^{(t)}}{\beta^{(t)}},$$

  where $w_{\text{cell}}^{(t)}$ is the normalized probability of a grid cell t, and $\alpha^{(t)}$ and $\beta^{(t)}$ are center locations of grid cells

- In 1D further variations, e.g. trapezoid

Quadrature integration

- The simplest quadrature integration is grid integration
  - Evaluate function in a grid and compute

$$E[-\alpha/\beta] \approx \sum_{t=1}^{T} w_{\text{cell}}^{(t)} \frac{\alpha^{(t)}}{\beta^{(t)}},$$

where $w_{\text{cell}}^{(t)}$ is the normalized probability of a grid cell t, and $\alpha^{(t)}$ and $\beta^{(t)}$ are center locations of grid cells

- In 1D further variations, e.g. trapezoid

- In 2D and higher
  - nested quadrature
  - product rules

Monte Carlo - history

- Used already before computers
    - Buffon (18th century; needles)
    - De Forest, Darwin, Galton (19th century)
    - Pearson (19th century; roulette)
    - Gosset (Student, 1908; hat)

Monte Carlo - history

- Used already before computers
    - Buffon (18th century; needles)
    - De Forest, Darwin, Galton (19th century)
    - Pearson (19th century; roulette)
    - Gosset (Student, 1908; hat)
- "Monte Carlo method" term was proposed by Metropolis, von Neumann or Ulam in the end of 1940s
    - they worked together in atomic bomb project
    - Metropolis and Ulam, "The Monte Carlo Method", 1949

Monte Carlo - history

- Used already before computers
  - Buffon (18th century; needles)
  - De Forest, Darwin, Galton (19th century)
  - Pearson (19th century; roulette)
  - Gosset (Student, 1908; hat)
- "Monte Carlo method" term was proposed by Metropolis, von Neumann or Ulam in the end of 1940s
  - they worked together in atomic bomb project
  - Metropolis and Ulam, "The Monte Carlo Method", 1949
- Bayesians started to have enough cheap computation time in 1990s
  - BUGS project started 1989 (last OpenBUGS release 2014)
  - Gelfand & Smith, 1990
  - Stan initial release 2012

Monte Carlo

- Simulate samples from the target distribution
  - these samples can be treated as any observations
- Use these samples, for example,
  - to compute means, deviations, quantiles
  - to draw histograms
  - to marginalize
  - etc.

Monte Carlo vs. deterministic

- Monte Carlo = simulation methods
  - evaluation points are selected stochastically (randomly)
- Deterministic methods (e.g. grid)
  - evaluation points are selected by some deterministic rule

How many simulation samples are needed?

- If samples are independent
  - usual methods to estimate the uncertainty due to a finite number of observations
- Markov chain Monte Carlo produces dependent samples
  - requires additional work to estimate the effective number of samples

How many simulation samples are needed?

- Expectation of unknown quantity

$$\mathsf{E}(\theta) \approx \frac{1}{L} \sum_l \theta^{(l)}$$

if L is big and $\theta^{(l)}$ are independent, way may assume that the distribution of the expectation approaches normal distribution (see Ch 4) with variance σ_θ^2/L (asymptotic normality)

  - this variance is independent on dimensionality of θ

How many simulation samples are needed?

- Expectation of unknown quantity

$$E(\theta) \approx \frac{1}{L} \sum_l \theta^{(l)}$$

  if L is big and $\theta^{(l)}$ are independent, way may assume that the distribution of the expectation approaches normal distribution (see Ch 4) with variance σ_θ^2 / L (asymptotic normality)

  - this variance is independent on dimensionality of θ
  - total variance is sum of the epistemic unecrtainty in the posterior and the uncertainty due to using finite number of Monte Carlo samples

  $$\sigma_\theta^2 + \sigma_\theta^2 / L$$

How many simulation samples are needed?

- Expectation of unknown quantity

$$E(\theta) \approx \frac{1}{L} \sum_l \theta^{(l)}$$

if L is big and $\theta^{(l)}$ are independent, way may assume that the distribution of the expectation approaches normal distribution (see Ch 4) with variance σ_θ^2/L (asymptotic normality)

- this variance is independent on dimensionality of θ
- total variance is sum of the epistemic unecrtainty in the posterior and the uncertainty due to using finite number of Monte Carlo samples

$$\sigma_\theta^2 + \sigma_\theta^2/L = \sigma_\theta^2(1 + 1/L)$$

How many simulation samples are needed?

- Expectation of unknown quantity

$$E(\theta) \approx \frac{1}{L} \sum_l \theta^{(l)}$$

if L is big and $\theta^{(l)}$ are independent, way may assume that the distribution of the expectation approaches normal distribution (see Ch 4) with variance σ_θ^2 / L (asymptotic normality)

- this variance is independent on dimensionality of θ
- total variance is sum of the epistemic unecrtainty in the posterior and the uncertainty due to using finite number of Monte Carlo samples

$$\sigma_\theta^2 + \sigma_\theta^2 / L = \sigma_\theta^2 (1 + 1/L)$$

- e.g. if $L = 100$, deviation increases by $\sqrt{1 + 1/L} = 1.005$ ie. Monte Carlo error is very small (for the expectation)

How many simulation samples are needed?

- Expectation of unknown quantity

$$E(\theta) \approx \frac{1}{L} \sum_l \theta^{(l)}$$

if L is big and $\theta^{(l)}$ are independent, way may assume that the distribution of the expectation approaches normal distribution (see Ch 4) with variance σ_θ^2/L (asymptotic normality)

- this variance is independent on dimensionality of θ
- total variance is sum of the epistemic unecrtainty in the posterior and the uncertainty due to using finite number of Monte Carlo samples

$$\sigma_\theta^2 + \sigma_\theta^2/L = \sigma_\theta^2(1 + 1/L)$$

- e.g. if $L = 100$, deviation increases by $\sqrt{1 + 1/L} = 1.005$ ie. Monte Carlo error is very small (for the expectation)
- See Ch 4 for counter-examples for asymptotic normality

Example: Kilpisjärvi summer temperature

Average temperature in June, July, and August at Kilpisjärvi, Finland

Example: Kilpisjärvi summer temperature

Average temperature in June, July, and August at Kilpisjärvi, Finland

Posterior fit with 90% interval

# Example: Kilpisjärvi summer temperature

Posterior of temperature change

C°/year

Example: Kilpisjärvi summer temperature

Posterior of temperature change

C°/century

Example: Kilpisjärvi summer temperature

Posterior of temperature change

$\sigma_\theta \approx 0.827$, MCSE ≈ 0.0827, total deviation ≈ 0.831

Example: Kilpisjärvi summer temperature

Posterior of temperature change

MCSE of posterior mean with S=1000

C°/century

$\sigma_\theta \approx 0.827$, MCSE ≈ 0.0261, total deviation ≈ 0.827

Example: Kilpisjärvi summer temperature

Posterior of temperature change

MCSE of posterior mean with S=100

0   2   4
C°/century

Posterior of temperature change

MCSE of posterior mean with S=1000

0   2   4
C°/century

Example: Kilpisjärvi summer temperature

Posterior of temperature change

MCSE of posterior mean with S=100

0 2 4
C°/century

Posterior of temperature change

MCSE of 5% quantile with S=100

0 2 4
C°/century

Posterior of temperature change

MCSE of posterior mean with S=1000

0 2 4
C°/century

Posterior of temperature change

MCSE of 5% quantile with S=1000

0 2 4
C°/century

Example: Kilpisjärvi summer temperature

Example: Kilpisjärvi summer temperature

Tail quantiles are more difficult to estimate

How many simulation samples are needed?

- Posterior probability

$$p(\theta \in A) \approx \frac{1}{L} \sum_l I(\theta^{(l)} \in A)$$

where $I(\theta^{(l)} \in A) = 1$ if $\theta^{(l)} \in A$

- $I(\cdot)$ is binomially distributed as $p(\theta \in A)$
  - $\rightarrow$ var$(I(\cdot)) = p(1-p)$ (Appendix A, p. 579)
  - $\rightarrow$ standard deviation of p is $\sqrt{p(1-p)/L}$

How many simulation samples are needed?

- Posterior probability

$$p(\theta \in A) \approx \frac{1}{L} \sum_l I(\theta^{(l)} \in A)$$

where $I(\theta^{(l)} \in A) = 1$ if $\theta^{(l)} \in A$

- $I(\cdot)$ is binomially distributed as $p(\theta \in A)$
  - $\rightarrow$ var$(I(\cdot)) = p(1-p)$ (Appendix A, p. 579)
  - $\rightarrow$ standard deviation of p is $\sqrt{p(1-p)/L}$
- if $L = 100$ and $p \approx 0.5$, $\sqrt{p(1-p)/L} = 0.05$
  ie. accuracy is about 5% units

How many simulation samples are needed?

- Posterior probability

$$p(\theta \in A) \approx \frac{1}{L} \sum_l I(\theta^{(l)} \in A)$$

where $I(\theta^{(l)} \in A) = 1$ if $\theta^{(l)} \in A$

- $I(\cdot)$ is binomially distributed as $p(\theta \in A)$
  - $\rightarrow$ var$(I(\cdot)) = p(1-p)$ (Appendix A, p. 579)
  - $\rightarrow$ standard deviation of p is $\sqrt{p(1-p)/L}$
- if $L = 100$ and $p \approx 0.5$, $\sqrt{p(1-p)/L} = 0.05$
  ie. accuracy is about 5% units
- $L = 2500$ simulation samples needed for 1% unit accuracy

How many simulation samples are needed?

- Posterior probability

$$p(\theta \in A) \approx \frac{1}{L} \sum_l I(\theta^{(l)} \in A)$$

where $I(\theta^{(l)} \in A) = 1$ if $\theta^{(l)} \in A$

- $I(\cdot)$ is binomially distributed as $p(\theta \in A)$
  - $\rightarrow$ var$(I(\cdot)) = p(1-p)$ (Appendix A, p. 579)
  - $\rightarrow$ standard deviation of p is $\sqrt{p(1-p)/L}$
- if $L = 100$ and $p \approx 0.5$, $\sqrt{p(1-p)/L} = 0.05$
  ie. accuracy is about 5% units
- $L = 2500$ simulation samples needed for 1% unit accuracy

- To estimate small probabilities, a large number of samples is needed
  - to be able to estimate p, need to get samples with $\theta^{(l)} \in A$, which in expectation requires $L \gg 1/p$

Example: Kilpisjärvi summer temperature

Posterior uncertainty p(temperature change < 0)

Example: Kilpisjärvi summer temperature

Posterior uncertainty p(temperature change < 0)

Example: Kilpisjärvi summer temperature

Posterior uncertainty p(temperature change < 0)

Uncertainty given S=100, and 0 draws < 0

Uncertainty given S=100, and 1 draw < 0
Uncertainty given S=100, and 2 draws < 0

p(temperature change < 0)

Example: Kilpisjärvi summer temperature

Posterior uncertainty p(temperature change < 0)

Uncertainty given S=1000, and 8 draws < 0

p(temperature change < 0)

Example: Kilpisjärvi summer temperature

Posterior uncertainty p(temperature change < 0)

Uncertainty given S=4000, and 34 draws < 0
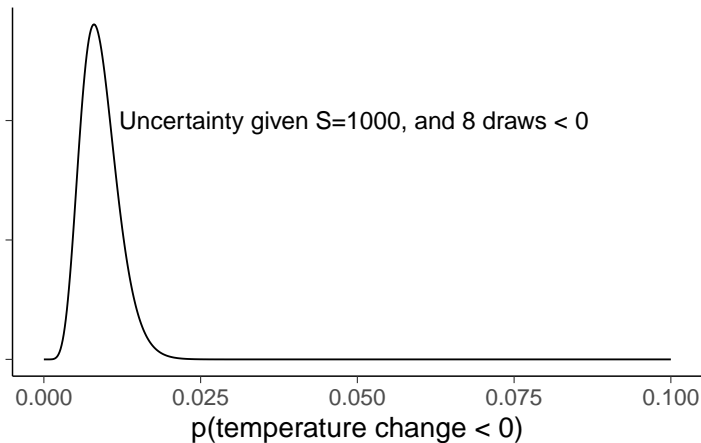
p(temperature change < 0)

How many simulation samples are needed?

- Less samples needed with
    - deterministic methods
    - marginalization (Rao-Blackwellization)
    - variance reduction methods, such, control variates

How many simulation samples are needed?

- Number of independent samples needed doesn't depend on the number of dimensions
  - but it may be difficult to obtain independent samples in high dimensional case

Direct simulation

- Produces independent samples
    - Using analytic transformations of unifrom random numbers (eg. appendix A)
    - factorization
    - numerical inverse-cdf
- Problem: restricted to limited set of models

Random number generators

- Good pseudo random number generators are sufficient for Bayesian inference
  - pseudo random generator uses deterministic algorithm to produce a sequence which is difficult to make difference from truly random sequence
  - modern software used for statistical analysis have good pseudo RNGs

Direct simulation: Example

- Box-Muller -method:
  If U_1 and U_2 are independent draws from distribution $U(0,1)$, and

$$X_1 = \sqrt{-2\log(U_1)}\cos(2\pi U_2)$$
$$X_2 = \sqrt{-2\log(U_1)}\sin(2\pi U_2)$$

  then X_1 and X_2 are independent samples from the distribution $N(0,1)$

Direct simulation: Example

- Box-Muller -method:
  If U_1 and U_2 are independent draws from distribution
  $U(0, 1)$, and

  $$X_1 = \sqrt{-2\log(U_1)}\cos(2\pi U_2)$$
  $$X_2 = \sqrt{-2\log(U_1)}\sin(2\pi U_2)$$

  then X_1 and X_2 are independent samples from the
  distribution $N(0, 1)$
    - not the fastest method due to trigonometric computations
    - for normal distrbution more than ten different methods
    - e.g. R uses inverse-CDF

Grid sampling and curse of dimensionality

- 10 parameters
- if we don't know beforehand where the posterior mass is
    - need to choose wide box for the grid
    - need to have enough grid points to get some of them where essential mass is
- e.g. 50 or 1000 grid points per dimension
    - $\rightarrow$ $50^{10} \approx$ 1e17 grid points
    - $\rightarrow$ $1000^{10} \approx$ 1e30 grid points
- R and my current laptop can compute density of normal distribution about 20 million times per second
    - $\rightarrow$ evaluation in 1e17 grid points would take 150 years
    - $\rightarrow$ evaluation in 1e30 grid points would take 1 500 billion years

Indirect sampling

- Rejection sampling
- Importance sampling
- Markov chain Monte Carlo (next week)

Rejection sampling

- Proposal forms envelope over the target distribution $q(\theta|y)/Mg(\theta) \leq 1$
- Draw from the proposal and accept with probability $Mg(\theta)/q(\theta|y)$

Rejection sampling

- Proposal forms envelope over the target distribution $q(\theta|y)/Mg(\theta) \leq 1$
- Draw from the proposal and accept with probability $Mg(\theta)/q(\theta|y)$

- Accepted - Rejected - - Mg(theta) — q(theta|y)

Rejection sampling

- Proposal forms envelope over the target distribution $q(\theta|y)/Mg(\theta) \leq 1$
- Draw from the proposal and accept with probability $Mg(\theta)/q(\theta|y)$
- Common for truncated distributions

- Accepted - Rejected - - Mg(theta) — q(theta|y)

Rejection sampling

- The number of accepted samples is the effective sample size
  - with bad proposal distribution may require a lot of trials
  - selection of good proposal gets very difficult when the number of dimensions increase
  - reliable diagnostics and thus can be a useful part

Importance sampling

- Proposal does not need to have a higher value everywhere

Target, proposal, and draws

Importance sampling

- Proposal does not need to have a higher value everywhere

Target, proposal, and draws

$$E[f(\theta)] \approx \frac{\sum_s w_s f(\theta^{(s)})}{\sum_s w_s}, \quad \text{where} \quad w_s = \frac{q(\theta^{(s)})}{g(\theta^{(s)})}$$

Importance sampling

- Proposal does not need to have a higher value everywhere

Draws and importance weights

$$E[f(\theta)] \approx \frac{\sum_s w_s f(\theta^{(s)})}{\sum_s w_s}, \quad \text{where} \quad w_s = \frac{q(\theta^{(s)})}{g(\theta^{(s)})}$$

Importance sampling

- Resampling using normalized importance weights can be used to pick a smaller number of draws with uniform weights
- Selection of good proposal gets more difficult when the number of dimensions increase
- Often used to correct distributional approximations

Importance sampling

- Variation of the weights affect the effective sample size
  - if single weight dominates, we have effectively one sample
  - if weights are equal, we have effectively S samples
- Central limit theorem holds only if variance of the weight distribution is finite
- See Vehtari, Gelman and Gabry (2017). Pareto smoothed importance sampling. arXiv preprint arXiv:1507.02646, https://arxiv.org/abs/1507.02646 for improved diagnostics and stability.
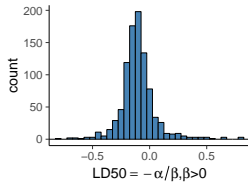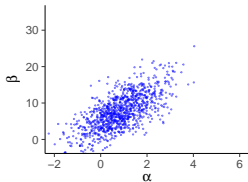
Exmple: Importance sampling in Bioassay
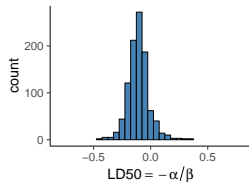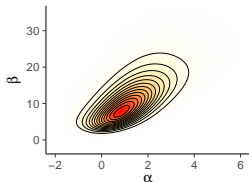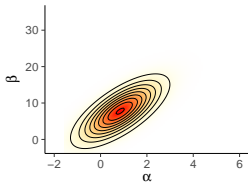
Exmple: Importance sampling in Bioassay
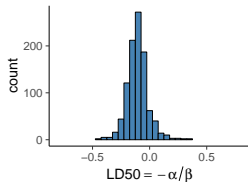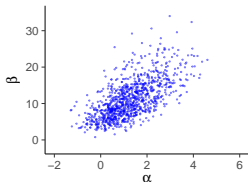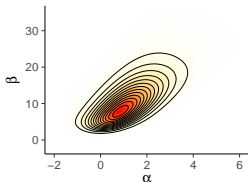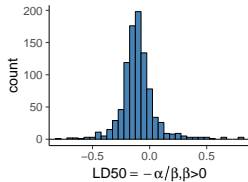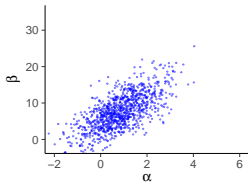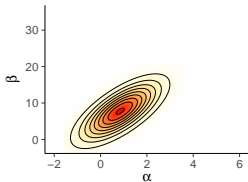
Normal approximation is discussed more in BDA3 Ch 4

Exmple: Importance sampling in Bioassay

Normal approximation is discussed more in BDA3 Ch 4
But the normal approximation is not that good here:
Grid sd(LD50) $\approx$ 0.1, Normal sd(LD50) $\approx$ .75!

# Exmple: Importance sampling in Bioassay

Exmple: Importance sampling in Bioassay
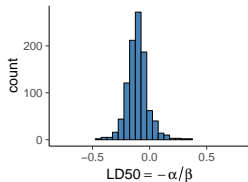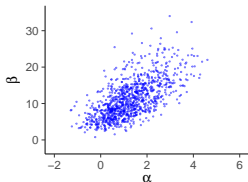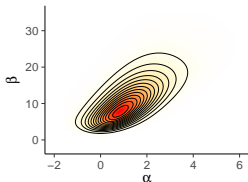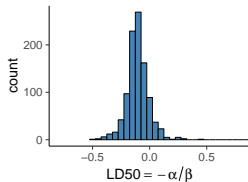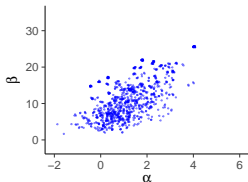
Grid sd(LD50) $\approx$ 0.1, SIR sd(LD50) $\approx$ 0.1

Exmple: Importance sampling in Bioassay

Exmple: Importance sampling in Bioassay

SIR

Exmple: Importance sampling in Bioassay

IS weights

Exmple: Importance sampling in Bioassay

IS weights

Exmple: Importance sampling in Bioassay

Exmple: Importance sampling in Bioassay

$$S_{\text{eff}} = \frac{1}{\sum_{s=1}^{S}(\tilde{w}(\theta^s))^2}, \quad \text{where } \tilde{w}(\theta^s) = w(\theta^s)/\sum_{s'=1}^{S} w(\theta^{s'})$$

Exmple: Importance sampling in Bioassay

IS weights

$$S_{\text{eff}} = \frac{1}{\sum_{s=1}^{S}(\tilde{w}(\theta^s))^2}, \quad \text{where } \tilde{w}(\theta^s) = w(\theta^s)/\sum_{s'=1}^{S} w(\theta^{s'})$$

$$S_{\text{eff}} \approx 270$$

Exmple: Importance sampling in Bioassay

IS weights

$$S_{\text{eff}} = \frac{1}{\sum_{s=1}^{S} (\tilde{w}(\theta^s))^2}, \quad \text{is based on variance of } \tilde{w}(\theta^s)$$

$$S_{\text{eff}} \approx 270$$

Exmple: Importance sampling in Bioassay

IS weights

$$S_{\text{eff}} = \frac{1}{\sum_{s=1}^{S}(\tilde{w}(\theta^s))^2}, \quad \text{is based on variance of } \tilde{w}(\theta^s)$$

$S_{\text{eff}} \approx 270$

Pareto-k diagnostic preferably < 0.7:

Exmple: Importance sampling in Bioassay

IS weights

$$S_{\mathrm{eff}} = \frac{1}{\sum_{s=1}^{S}(\tilde{w}(\theta^s))^2}, \quad \text{is based on variance of } \tilde{w}(\theta^s)$$

$$S_{\mathrm{eff}} \approx 270$$

Pareto-k diagnostic preferably < 0.7: $\hat{k} \approx 0.76$

Pareto smoothed importance sampling

- Pareto-k diagnostic estimate the number of existing moments ($\lfloor 1/k \rfloor$)
- Finite variance and central limit theorem for $k < 1/2$
- Finite mean and generalized central limit theorem for $k < 1$, but pre-asymptotic constant grows impractically large for $k > 0.7$
- See Vehtari, Gelman and Gabry (2017). Pareto smoothed importance sampling. arXiv preprint arXiv:1507.02646, https://arxiv.org/abs/1507.02646 for improved diagnostics and stability.

Importance sampling leave-one-out cross-validation

- Later in the course you will learn how $p(\theta|y)$ can be used as a proposal distribution for $p(\theta|y_{-i})$
  - which allows fast computation of leave-one-out cross-validation

$$p(y_i|y_{-i}) = \int p(y_i|\theta)p(\theta|y_{-i})d\theta$$

Curse of dimensionality

- Number of grid points increases exponentially
- Concentration of the measure, ie, where is the most of the mass?

Markov chain Monte Carlo (MCMC)

- Markov chain goes where most of the posterior mass is
- Certain MCMC methods scale well to high dimensions
- MCMC methods in this course
  - Gibbs: "iterative conditional sampling"
  - Metropolis: "random walk in joint distribution"
  - Dynamic Hamiltonian Monte Carlo: "state-of-the-art" used in Stan